




* a single frozen LLM plays all three roles — no fine-tuning and no external reward model