

FinOpsBench: A Benchmark for Evaluating Agentic LLMs on Financial Analysis Tasks with Tool Use

Anonymous ACL submission

Abstract

We introduce FinOpsBench, a comprehensive benchmark for evaluating the capabilities of large language models (LLMs) in agentic financial analysis tasks that require tool use. The benchmark is released in two versions: (1) FinOpsBench-v1, a synthetic set of 5,979 natural-language financial analysis tasks that an agent must solve by issuing structured queries to a tool-callable data environment, with **quality control by a panel of LLM judges**, and (2) FinOpsBench-v2, a set of 1,108 multi-hop reasoning tasks grounded in real financial documents and embedded in custom tool environments. Both subsets require models to retrieve information exclusively through tool calls rather than relying on in-context information, **testing genuine agentic tool use and reasoning capabilities**. We evaluate a range of frontier and open-source models, finding that even the strongest model achieves only 68.9% on FinOpsBench-v1 and 69.6% on FinOpsBench-v2, while open-source baselines fall as low as 21.9% and 16.3% respectively, demonstrating significant room for improvement in agentic financial reasoning. We release the full benchmark, evaluation code, and model traces to facilitate future research.

1 Introduction

Large language models (LLMs) are increasingly deployed as autonomous agents capable of using tools to accomplish complex tasks (Schick et al., 2023; Qin et al., 2024b). In the financial domain, such agents can assist with financial analysis, compliance checking, and decision support by issuing tool calls against structured business data and synthesizing the retrieved evidence into actionable answers. However, evaluating the agentic capabilities of LLMs presents significant challenges. Existing financial QA benchmarks like FinQA (Chen et al., 2021) and TAT-QA (Zhu et al., 2021) primarily test reading comprehension and numerical




Figure 1: Accuracy on FinOpsBench grows log-linearly with the (estimated) parameter count of the agent’s base model. Closed-source sizes are taken from the IKP estimates (Li, 2026)

reasoning over provided contexts, rather than genuine tool-use and multi-step reasoning. We introduce FinOpsBench, a benchmark specifically designed to evaluate agentic LLMs on financial analysis tasks requiring tool use. **We argue an agentic financial benchmark should at once provide diagnostic control over a large pool of executable tasks (so that failures can be traced to a specific planning or tool-use mistake rather than to evaluator noise) and realism on human-authored financial questions (so that accuracies remain comparable with the financial-QA literature).** These two goals pull in opposite directions, so FinOpsBench is released as two complementary parts - each targeting one of these goals and both evaluated under the same agent environment:

- **FinOpsBench-v1:** 5,979 examples; each pairs a natural-language financial question with a small, freshly generated database that the agent must query through a single structured-data tool to reach the answer. Examples are quality-controlled by a panel of LLM judges.
- **FinOpsBench-v2:** This benchmark is pro-

duced by building agentic environments around questions and tables taken from FinQA, and consists of 1,108 samples. Each sample bundles a financial question, a tool environment with diverse custom Python tools, distractor tools and rows, and a gold answer. The agent must perform multi-hop reasoning across several tool calls.

Both benchmarks share the same design principles. First, answers cannot be extracted from the prompt alone, so agents must use tools to retrieve information. Second, examples require multi-step reasoning with intermediate tool calls. Third, distractor information is present to test the agent’s ability to filter relevant from irrelevant data. A related concern with benchmarks is data contamination, where test examples may have been seen during model training. Where part of our benchmark builds on a prior dataset, each item is re-instantiated inside a bespoke tool environment with a redacted system prompt and a multi-step tool-call plan the model has not seen during training, so familiarity with the underlying source items does not translate into a direct answer pathway.

Our main contributions are:

- A novel benchmark for evaluating agentic LLMs on financial tasks requiring genuine tool use.
- A rigorous benchmark construction pipeline with quality control by a panel of LLM judges
- Comprehensive evaluation of a range of frontier and open-source LLMs used as the *base models* inside the agents, revealing significant gaps in agentic financial reasoning capabilities.
- An open-source release¹ of the complete benchmark together with the evaluation framework and all model traces. We treat this release as a contribution in its own right, supporting reproducible research and direct reuse of our infrastructure for new agentic financial tasks.

As Figure 1 shows that the performance of different models on FinOpsBench grows log-linearly with the parameter count of the agent’s base model. This is an expected sanity check that the benchmark’s difficulty tracks model capability rather than a task-design artifact, and that performance gaps therefore reflect real differences in agentic competence. Parameter counts for closed-source models

are taken from the IKP calibration of Li (2026).

2 Related Work

FinOpsBench is situated at the intersection of three research areas: context-bound financial Question Answering (QA), general agentic tool-use evaluation, and domain-specific semantic parsing.

Context-Bound Financial QA. Foundational financial benchmarks like FinQA (Chen et al., 2021), ConvFinQA (Chen et al., 2022), TAT-QA (Zhu et al., 2021), MultiHiertt (Zhao et al., 2022), and FinanceBench (Islam et al., 2023), together with the broader PIXIU benchmark suite (Xie et al., 2023), measure numerical reasoning, tabular-textual integration, and reasoning over SEC filings within static, explicitly provided contexts. These established datasets are inherently non-agentic as they do not require dynamic interaction or external system calls; all necessary information is supplied in the prompt, testing reading comprehension rather than tool use or dynamic information retrieval. Our FinOpsBench-v2 component transitions this evaluation to a dynamic environment where information must be actively retrieved through multi-hop tool calls, isolating and testing the core planning and execution logic.

Agentic and Tool-Use Benchmarks. General agent benchmarks, such as ToolBench (Qin et al., 2024a,b), MCP-Verse (Lei et al., 2025), AgentBench (Liu et al., 2024), and SWE-bench (Jimenez et al., 2024), establish the methodological standard for evaluating LLMs’ planning abilities against a large breadth of tools and domains (web browsing, code execution, APIs at scale). While essential for assessing general tool augmentation, these benchmarks prioritize scale over domain depth, and none concentrate on financial analysis with structured-data tools. The wider agentic-evaluation landscape spans web navigation (WebArena (Zhou et al., 2024), WebShop (Yao et al., 2022)), embodied text-world tasks (ALFWorld (Shridhar et al., 2021)), general AI assistance over the open web (GAIA (Mialon et al., 2024)), retail and airline tool-use dialogues (τ -bench (Yao et al., 2024)), and tool-selection decisions (MetaTool (Huang et al., 2024)); none of these target financial-analysis tasks either. Financial analysis critically demands precision and specialized interpretation. By specializing the tool environment for finance, FinOpsBench tests the *depth* of financial logic, revealing significant gaps where general planning skills fail to confer

¹Code: <https://anonymous.4open.science/r/FinOpsBench>

the necessary domain competence (16–70% accuracy on grounded financial reasoning tasks). Recent real-world financial benchmarks like FinGAIA (Zeng et al., 2025), the Finance Agent Benchmark (Bigeard et al., 2025), and FinAgentBench (Choi et al., 2025) prioritize realism by using live web data (SEC EDGAR) or SEC-filing retrieval, while FinOpsBench is engineered for *diagnostic fidelity* using controlled, synthetic environments to attribute failures unequivocally to flawed planning logic, rather than external noise or API variability. We further compare FinOpsBench against FinAgentBench and the Finance Agent Benchmark numerically along the size–accuracy axis in Appendix A.

Tool Use over Structured Business Data A related line of work studies semantic parsing over structured data: cross-domain datasets such as Spider (Yu et al., 2018) and BIRD (Li et al., 2023) score a model on the fidelity of a single generated query string against a fixed schema. FinOpsBench is deliberately *not* a semantic-parsing benchmark: the agent is never scored on a query string. Instead, FinOpsBench-v1 evaluates the full agentic loop (tool selection, multi-turn interaction with the data tool, evidence aggregation across calls, and natural-language answer synthesis) over financially specialized environments. Two consequences follow that make the task qualitatively different from text-to-query parsing. First, success requires planning under partial observability: the schema and contents must be probed via tool calls rather than read from context. Second, financial domain semantics such as accounts-payable aging, variance attribution, and revenue recognition become first-class evaluation targets rather than incidental surface features. The structured-data tool is the means, not the end being measured.

Our approach leverages synthetic data generation and a rigorous quality control process, utilizing a panel of three LLM judges (Zheng et al., 2023; Verga et al., 2024), to mitigate data contamination risks and ensure high integrity at scale, validating that observed low accuracy reflects genuine performance limitations.

Data Generation. Synthetic data generation has proven valuable for creating evaluation benchmarks at scale (Wang et al., 2023). Recent work has explored using LLMs to generate and verify synthetic datasets with quality control mechanisms (Prabhakar et al., 2026). Our approach utilizes these methods and adapts them to creating financial agen-

Property	FinOpsBench-v1	FinOpsBench-v2
# Examples	5,979	1,108
Tool type	Single structured-data tool	Many custom Python tools
Multi-hop	≥ 1 hop	≥ 2 hops
Distractor data	Yes	Yes
Quality control	Panel	Execution-based

Table 1: Overview of the two FinOpsBench versions.

tic datasets.

3 Benchmark Description

3.1 Overview

Following the rationale laid out in the introduction, FinOpsBench comprises two complementary parts (Table 1): FinOpsBench-v1 materialises the diagnostic-control side of the split (a large, fully synthetic pool with executable ground truth), while FinOpsBench-v2 materialises the realism side, anchored in human-authored questions over real financial documents.

3.2 FinOpsBench-v1

FinOpsBench-v1 consists of financial analysis tasks that an agent must solve by issuing structured-data queries through a single tool. Each example is backed by an isolated SQLite store generated specifically for that example; the agent is never required to produce or score a query string, and only the final answer is evaluated. Each example contains *User role* (the persona making the query, e.g., CFO, auditor, financial analyst), *Query* (a natural language question about financial data), *Backing data schema* (a plausible company-data schema generated for this example), *Backing data* (realistic financial records including distractor rows), *Expected answer* (the ground-truth answer), and *Agent trace* (a reference solution demonstrating correct tool use).

FinOpsBench-v1 spans diverse financial analysis categories including Accounts Payable (AP) analysis, Revenue recognition, Variance analysis, Data integrity and reconciliation, and Financial reporting. Representative queries spanning user roles and task categories are listed in Appendix D, and the empirical category mix of the full FinOpsBench-v1 stream is reported in Appendix G.

3.3 FinOpsBench-v2

FinOpsBench-v2 transforms FinQA samples into complete agentic environments. Each example in-

cludes *System prompt* (scenario description and available tools without table data), *Synthetic backing store* (an SQLite (Owens and Allen, 2010) table set reconstructed from the original FinQA tables), *Core tools* (functions required to derive the answer), *Distractor tools* (irrelevant tools to test filtering ability), and *Correct plan* (reference multi-hop solution). Critically, the agent receives only the scenario description and tool definitions, not the underlying data table. This ensures agents must use tools to retrieve information rather than extracting answers from context. Representative questions, together with a worked example showing the query, available tools, and a reference tool-call sequence, are given in Appendix E and Table 5.

4 Benchmark Construction

4.1 FinOpsBench-v1 pipeline

FinOpsBench-v1 is created through a 9-stage pipeline. Figure 2 (top) gives an end-to-end overview of the FinOpsBench-v1 construction process; the individual stages are described in detail below.

Stage 1: Query Generation. We start with a few seed queries spanning financial analysis categories (in our experiment, we took 12 seed queries). These seeds are then expanded using an LLM, varying user roles and specific details. We utilize a prompt that takes one seed and generates 20 queries to increase diversity and run it multiple times to generate a large number of queries. We also use embedding-based similarity filtering (cosine threshold 0.9) to remove near-duplicates.

Stage 2: Schema Generation. For each query, an LLM generates a plausible data schema that is: (1) sufficient to answer the query, (2) realistic for a company’s records, and (3) appropriately complex.

Stage 3: Data Generation. The LLM is then tasked with generating realistic data for each query and schema, including: (1) data necessary to answer the query, and (2) distractor rows to increase difficulty. The LLM is also required to provide the correct answer to the query consistent with the generated data.

Stage 4: Execution-Based Validation. All generated schema and data definitions are validated by execution against SQLite. Errors are iteratively fixed using LLM assistance (up to 5 attempts). We found that LLMs struggle with fixing small errors in (relatively) long generated programs, so we split the program into individual statements and fix them

one by one.

Stage 5: Agent Trace Generation. An agent runs on each query. It is supplied with access to the example’s backing data through the structured-data tool, with up to 6 rounds of interaction allowed. The full conversation trace is recorded.

At this point, we have a benchmark subset that contains a query, a data schema, backing data (records necessary to answer the query and distractor rows), intended answer to the query, and an agent trace demonstrating the solution to the query. However, there may be hallucinations, unnatural data, wrong tool calls, or wrong answers in the agent trace that need to be identified and fixed.

Stage 6: Panel Judgement. A panel of three LLM judges evaluates each example’s data and agent trace on five criteria: (1) *Data naturalness*: Is the generated data realistic? (2) *Trace reasonableness*: Does the agent examine appropriate data? (3) *Trace soundness*: Is the reasoning logical? (4) *Grounding*: Is reasoning based only on tool outputs? (5) *Answer correctness*: Does the answer address the query?

An example passes only if it receives majority approval (2/3 judges) on *all* criteria. In that case, it proceeds to the next stage. Otherwise, it goes to the improvement loop (Stages 7 and 8). Note that the loop only runs once for each example; if the improved example fails the judgement again, it is rejected completely.

Stage 7: Feedback Reconciliation. The judges’ criticisms are given to an LLM which is asked to aggregate them into actionable feedback. The LLM at this stage has access to most of the example’s data, including the query, the data schema, the backing data, and the agent trace.

Stage 8: Feedback Application. The agent runs again with a prompt containing the query, the data schema, the previous agent trace, and the feedback produced in Stage 7. It still has access to the structured-data tool, and the conversation is limited to 6 rounds. The full conversation trace is recorded.

Stage 9: Second Judgement. The improved examples are then re-judged by the panel. If an example passes, it proceeds to the next stage. Otherwise, it is dropped from the benchmark.

Final Filtering. The final benchmark includes only examples that passed panel judgement. We run two filtering passes to ensure quality and complexity of the benchmark. The first pass uses an LLM to compare the final answer produced in the agent trace with the intended answer created in Stage

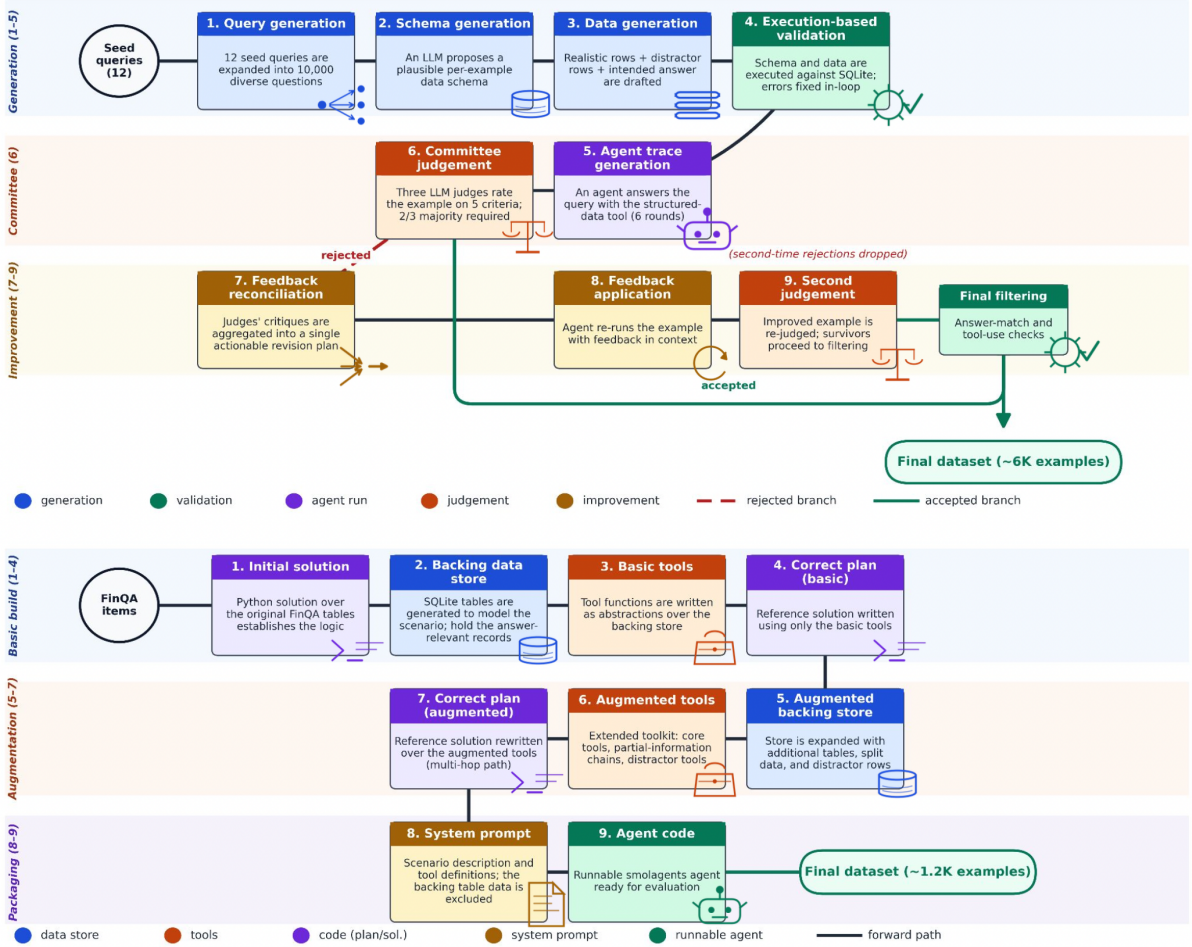


Figure 2: Construction pipelines for the two FinOpsBench versions. **Top:** FinOpsBench-v1 (9 stages with a panel-judgement gate and an improvement loop). **Bottom:** FinOpsBench-v2 (9 strictly linear stages, each validated by execution).

3. Additionally, we filter out examples where the agent did not actually **invoke the structured-data tool** (removing examples answerable from context alone).

4.2 FinOpsBench-v2 pipeline

The FinOpsBench-v2 pipeline transforms FinQA samples through 9 stages (Figure 2, bottom):

Stage 1: Initial Solution. An LLM creates a direct Python solution using the original table data, establishing the computational logic.

Stage 2: Backing Data Store Generation. An SQLite-backed table set is generated to model the scenario, containing the records needed to answer the question.

Stage 3: Basic Tools. Tool functions are created as abstractions over the **backing store**, designed to enable multi-hop reasoning.

Stage 4: Correct Plan (Basic). A reference solution is written using the basic tools.

Stage 5: Augmented Backing Store. The **backing store** is expanded with additional tables, split data, and distractor rows to increase difficulty.

Stage 6: Augmented Tools. Extended tools are created including core tools for the correct solution, partial information tools requiring chaining, and distractor tools providing irrelevant data.

Stage 7: Correct Plan (Augmented). A reference solution using the augmented tools demonstrates the multi-hop solution path.

Stage 8: System Prompt. A system prompt is generated containing the scenario description and tool definitions, explicitly excluding table data and answer hints.

Stage 9: Agent Code. A runnable agent using the `smolagents` (Roucher et al., 2025) library is generated for evaluation.

Each stage is required to output a file (Python code in all stages except Stage 8 which outputs a system prompt with the task formulation). For each




Figure 3: FinOpsBench-v1 construction funnel: number of examples surviving each stage of the 9-stage pipeline plus the final filtering pass.

stage, the LLM receives in the context: (1) A long description of the whole process with all the stages described; (2) An indication of what stage is currently running; (3) All previous files’ content from the previous stages, and (4) In Stage 9, an example agent code showing coding practices for achieving more stable results. After completing, each stage undergoes an execution-based validation. If the execution fails, the stage is rerun with a prompt that includes the error message and asks to fix the errors (up to 10 attempts).

5 Experiments

5.1 Setup

We evaluate a diverse set of LLMs as the agent’s base model. The frontier set is GPT-4.1, GPT-4.1-mini, GPT-5, GPT-5-mini, and o4-mini (Achiam et al., 2023) (OpenAI). The open-source set is Qwen3-8B and Qwen3-30B-A3B (Yang et al., 2025) (Alibaba), and Llama-3.1-8B-Instruct (Dubey et al., 2024) (Meta). Frontier models are accessed through OpenAI’s API; the three open-source models are served locally on a single NVIDIA H100 GPU. Our primary metric is *accuracy*. For FinOpsBench-v1, an LLM judge compares the agent’s final answer against the expected output and outputs a yes/no verdict; we evaluate on a randomly selected test set of 729 examples for cost reasons. For FinOpsBench-v2, the numerical answers are compared against the correct plan output with a tolerance of one least significant digit, and we evaluate on the full set.

5.2 Construction statistics

We first describe how many examples survived each pipeline stage for the two versions.

Stage	Number of Examples
Seed queries	12
1. Query expansion	10,000
2. Schema generation	10,000
3. Data generation	10,000
4. Execution-based validation	10,000
5. Agent trace generation	9,557
6. Panel judgement	
– passed	5,401
– to be improved	4,156
7. Feedback reconciliation	4,156
8. Feedback application	3,967
9. Panel judgement	2,832
Total passed so far	8,233
Filtering (post-pipeline)	5,979

Table 2: Number of examples in FinOpsBench-v1 after each stage.

FinOpsBench-v1. We constructed FinOpsBench-v1(5,979 examples) by starting with 12 seed queries and generating 10,000 diverse queries, then proceeding with all the remaining stages of the pipeline. The default LLM used in all stages is GPT-4.1-mini with the following exceptions: Stages 4 (execution-based validation), 7 (feedback reconciliation), and 8 (feedback application) are run with o4-mini; Stage 6 (panel judgement) is run with three different models: Claude Sonnet 4 (Anthropic, 2025), o4-mini, and o3-mini, one per judge. See Table 2 for the number of examples after each stage.

Figure 3 visualises the same attrition: the pipeline holds steady through Stages 1–4 (deterministic generation against a fixed schema), then loses 4% of examples at the agent-trace stage, a further 14% at panel judgement, and 27% at the final filtering pass.

FinOpsBench-v2. We ran the data-example creation routine using the protocol described above for 1,247 examples drawn uniformly at random from the FinQA training split, resulting in 1,108 examples (the rest failed in at least one of the stages). Appendix F lists the per-example file layout.

5.3 Results

The final evaluation results for both versions of the benchmark are presented in Table 3.

Findings. Having run experiments on both versions, we observe a consistent picture: overall model capability tracks benchmark accuracy, but the strongest models still leave substantial headroom. The log-linear dependence of accuracy

Model	Acc. v1 (%)	Acc. v2 (%)
<i>Frontier models</i>		
GPT-5	68.9	69.6
GPT-5-mini	65.8	67.5
o4-mini	67.1	67.3
GPT-4.1	62.4	60.6
GPT-4.1-mini	61.5	56.9
<i>Open-source models</i>		
Qwen3-30B-A3B	50.5	53.0
Qwen3-8B	47.6	44.1
Llama-3.1-8B-Instruct	21.9	16.3

Table 3: Per-model accuracy on FinOpsBench-v1(v1) and FinOpsBench-v2(v2). Bold values mark the best score within each tier.

on the agent’s base-model parameter count is already visible in Figure 1, and Table 3 provides the per-model numbers. On FinOpsBench-v1, even the best model (GPT-5) reaches only 68.9%, indicating substantial room for improvement; modern and larger models such as GPT-5 outperform smaller, older ones such as Llama-3.1-8B; and the frontier/open-source gap is wide, with even Qwen3-30B-A3B scoring 11 pp below the worst frontier model (GPT-4.1-mini). Figure 5 further shows that this ranking is stable across the four query categories carried by the seed set: every model’s per-category polygon is approximately concentric, so the gap is not driven by any single financial sub-domain.

On FinOpsBench-v2 we observe essentially the same ranking. GPT-5 and o4-mini lead at ~67–70%, the GPT-4.1 variants follow, and the open-source models trail. Llama-3.1-8B again severely underperforms its size-matched newer sibling Qwen3-8B, suggesting that tool-use capability is sensitive to the training paradigm and not only to parameter count. Per-model accuracies track each other across the two versions (mean absolute difference 2.6 pp); the side-by-side bar chart in Figure 4 and the $y = x$ scatter in Appendix B (Figure 9) make this agreement explicit, with the only material outliers being Qwen3-30B-A3B and Qwen3-8B, which both score noticeably lower on the FinQA-grounded version, where the multi-hop, multi-tool environment exposes more of the weaker models’ failure modes.

Native tool calling vs. ReAct. While analyzing the traces of smaller models, we found it hard to understand what they were struggling with when




Figure 4: Per-model accuracy on FinOpsBench-v1 and FinOpsBench-v2. Frontier models occupy the left half, open-source the right.




Figure 5: Per-category accuracy on FinOpsBench-v1 for the eight evaluated models. Axes are the five categories described in Section 3.2; each query in the 729-example test set is assigned to one of these five categories. The category mix of the full FinOpsBench-v1 stream is reported in Appendix G.

calling tools, because there was no record of the reasoning behind the tool and argument choices. We therefore added a second experiment using the ReAct tool-calling protocol (Yao et al., 2023), which requires the model to output its reasoning before calling a tool. This experiment was performed on FinOpsBench-v1; Table 4 presents the results.

The split is clean and visible in Figure 6: every *non-thinking* model gains from switching to ReAct (a strict green delta on the right of the plot), while every *thinking* model loses (a strict red delta on the left). The gains shrink as the non-thinking model gets stronger, from +6.4 pp for Llama-3.1-8B-Instruct down to +1.4 pp for GPT-4.1, suggesting that ReAct’s explicit reasoning scaffold matters most when the model lacks an internal chain-of-thought process (Wei et al., 2022) and matters less, or even hurts, once the model already reasons be-

Model	Accuracy (%)	
	Native	ReAct
<i>Non-thinking Models</i>		
GPT-4.1	62.4	<u>63.8</u>
GPT-4.1-mini	61.5	<u>63.5</u>
Llama-3.1-8B-Instruct	21.9	<u>28.3</u>
<i>Thinking Models</i>		
GPT-5	<u>68.9</u>	64.6
o4-mini	<u>67.1</u>	64.3
Qwen3-30B-A3B	<u>50.5</u>	46.2
Qwen3-8B	<u>47.6</u>	44.3

Table 4: Evaluation results on tool calling protocol. Underlined values are the best result for each model.




Figure 6: Native tool calling vs. ReAct on FinOpsBench-v1. Numbers above each pair give the ReAct minus Native delta in percentage points.

fore acting.

6 Discussion

Benchmarking Agentic Capabilities. Our results demonstrate that FinOpsBench effectively discriminates between models’ agentic capabilities. The significant gap from ceiling performance (around 69%) suggests current models have substantial room for improvement in financial tool use.

Open-Source Gap. The large gap between frontier and open-source models (up to 50+ percentage points) highlights an important direction for open-source development. Financial agentic tasks may require specialized training on tool-use trajectories. Concretely, closing this gap likely requires trajectory-level fine-tuning or reinforcement learning over realistic tool-use sequences (including tool selection, argument formation, and intermediate-result handling) rather than additional next-token pre-training on financial text alone.

Native Tool Calling vs. ReAct. ReAct allows the model to reason before choosing a tool to call. This

has a large impact on performance for models that do not reason in their default setting, highlighting the improvement that test-time scaling can bring. For models that are already thinking, ReAct decreases performance. This may reflect that modern thinking models are well-optimized for structured function calling, while ReAct’s verbose format introduces additional opportunities for error.

Quality Control. Our panel-based quality control ensures high-quality examples but introduces potential biases. The judges themselves are LLMs, which may share systematic blind spots with the evaluated models.

7 Conclusion

We introduced FinOpsBench, a benchmark for evaluating agentic LLMs on financial analysis tasks requiring tool use. The benchmark is released in two versions: FinOpsBench-v1 (5,979 synthetically generated examples) and FinOpsBench-v2 (1,108 FinQA-grounded examples). Evaluation of frontier and open-source LLMs used as the base models inside the agents reveals significant room for improvement, with best accuracies of 69%.

Our work provides a foundation for developing and evaluating financial AI agents. Future work may extend the benchmark to additional financial domains, incorporate real corporate data sources (with appropriate anonymization), and develop training approaches to improve agentic financial reasoning.

Ethics Statement

We plan to release FinOpsBench under the MIT License for research and educational use. FinOpsBench-v2 is built on top of FinQA (Chen et al., 2021), which is also distributed under the MIT License; our use of FinQA (building derivative tool environments around its public questions and tables) is consistent with its intended use as a public research benchmark, and we redistribute only derivative artifacts, not the original FinQA tables.

FinOpsBench contains no personal or identifying data: FinQA draws on aggregated public S&P 500 disclosures and FinOpsBench-v1 is generated end-to-end, so no additional anonymisation was required. All artifacts are provided in English.

During the preparation of our work, AI assistants were used for auxiliary tasks. A detailed descrip-

tion of this process is given in Appendix H.

Limitations

FinOpsBench is scoped accordingly. Both versions currently instrument structured-data and Python-tool environments; other modalities such as live web search, document retrieval, and multimodal financial reports lie outside this scope. These are deliberate scope choices for an initial release rather than fundamental limitations, and the construction methodology can in principle accommodate such extensions.

References

Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, and 1 others. 2023. GPT-4 technical report. *arXiv preprint arXiv:2303.08774*.

Anthropic. 2024. The claude 3 model family: Opus, sonnet, haiku. Technical report, Anthropic.

Anthropic. 2025. System card: Claude Opus 4 and Claude Sonnet 4. Technical report, Anthropic.

Antoine Bigeard, Langston Nashold, Rayan Krishnan, and Shirley Wu. 2025. Finance Agent Benchmark: Benchmarking LLMs on real-world financial research tasks. *arXiv preprint arXiv:2508.00828*.

Zhiyu Chen, Wenhu Chen, Charese Smiley, Sameena Shah, Iana Borova, Dylan Langdon, Reema Moussa, Matt Beane, Ting-Hao Huang, Bryan Routledge, and William Yang Wang. 2021. FinQA: A dataset of numerical reasoning over financial data. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 3697–3711. Association for Computational Linguistics.

Zhiyu Chen, Shiyang Li, Charese Smiley, Zhiqiang Ma, Sameena Shah, and William Yang Wang. 2022. ConFinQA: Exploring the chain of numerical reasoning in conversational finance question answering. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 6279–6292. Association for Computational Linguistics.

Chanyeol Choi, Jihoon Kwon, Alejandro Lopez-Lira, Chaewoon Kim, Minjae Kim, Juneha Hwang, Jaeseon Ha, Hojun Choi, Suyeol Yun, Yongjin Kim, and 1 others. 2025. FinAgentBench: A benchmark dataset for agentic retrieval in financial question answering. In *Proceedings of the 6th ACM International Conference on AI in Finance (ICAIF)*, pages 632–637. Association for Computing Machinery.

Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela

Fan, and 1 others. 2024. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.

Yue Huang, Jiawen Shi, Yuan Li, Chenrui Fan, Siyuan Wu, Qihui Zhang, Yixin Liu, Pan Zhou, Yao Wan, Neil Zhenqiang Gong, and Lichao Sun. 2024. Meta-Tool benchmark for large language models: Deciding whether to use tools and which to use. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*.

Pranab Islam, Anand Kannappan, Douwe Kiela, Rebecca Qian, Nino Scherrer, and Bertie Vidgen. 2023. FinanceBench: A new benchmark for financial question answering. *arXiv preprint arXiv:2311.11944*.

Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik R. Narasimhan. 2024. SWE-bench: Can language models resolve real-world GitHub issues? In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*.

Fei Lei, Yibo Yang, Wenxiu Sun, and Dahua Lin. 2025. MCPVerse: An expansive, real-world benchmark for agentic tool use. *arXiv preprint arXiv:2508.16260*.

Bojie Li. 2026. Incompressible knowledge probes: Estimating black-box LLM parameter counts via factual capacity. *arXiv preprint arXiv:2604.24827*.

Jinyang Li, Binyuan Hui, Ge Qu, Jiaxi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Ruiying Geng, Nan Huo, Xuanhe Zhou, Chenhao Ma, Guoliang Li, Kevin C.-C. Chang, Fei Huang, Reynold Cheng, and Yongbin Li. 2023. Can LLM already serve as a database interface? a big bench for large-scale database grounded text-to-SQLs. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, volume 36.

Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su, Huan Sun, and 3 others. 2024. AgentBench: Evaluating LLMs as agents. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*.

Leland McInnes, John Healy, Nathaniel Saul, and Lukas Großberger. 2018. UMAP: Uniform manifold approximation and projection. *Journal of Open Source Software*, 3(29):861.

Meta AI. 2025. The Llama 4 herd: The beginning of a new era of natively multimodal AI innovation. Technical report, Meta.

Grégoire Mialon, Clémentine Fourrier, Thomas Wolf, Yann LeCun, and Thomas Scialom. 2024. GAIA: A benchmark for general AI assistants. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*.

Mike Owens and Grant Allen. 2010. *SQLite*. Apress LP New York.

700	Akshara Prabhakar, Zuxin Liu, Ming Zhu, Jianguo	Qianqian Xie, Weiguang Han, Xiao Zhang, Yanzhao	757
701	Zhang, Tulika Manoj Awalganekar, Shiyu Wang,	Lai, Min Peng, Alejandro Lopez-Lira, and Jimin	758
702	Zhiwei Liu, Haolin Chen, Thai Hoang, Juan Car-	Huang. 2023. PIXIU: A comprehensive benchmark,	759
703	los Niebles, and 1 others. 2026. Apigen-mt: Agentic	instruction dataset and large language model for fi-	760
704	pipeline for multi-turn data generation via simulated	nance. In <i>Advances in Neural Information Process-</i>	761
705	agent-human interplay. <i>Advances in Neural Informa-</i>	<i>ing Systems (NeurIPS), Datasets and Benchmarks</i>	762
706	<i>tation Processing Systems</i> , 38.	<i>Track</i> , volume 36.	763
707	Yujia Qin, Shengding Hu, Yankai Lin, Weize Chen,	An Yang, Baosong Anil, and 1 others. 2025. Qwen3	764
708	Ning Ding, Ganqu Cui, Zheni Zeng, Xuanhe Zhou,	technical report. <i>arXiv preprint arXiv:2505.09388</i> .	765
709	Yufei Huang, Chaojun Xiao, and 1 others. 2024a.	Shunyu Yao, Howard Chen, John Yang, and Karthik	766
710	Tool learning with foundation models. <i>ACM Com-</i>	Narasimhan. 2022. WebShop: Towards scalable real-	767
711	<i>puting Surveys</i> , 57(4):1–40.	world web interaction with grounded language agents.	768
712	Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan	In <i>Advances in Neural Information Processing Sys-</i>	769
713	Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang,	<i>tems (NeurIPS)</i> , volume 35.	770
714	Bill Qian, Sihan Zhao, Runchu Tian, Ruobing Xie,	Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik	771
715	Jie Zhou, Mark Gerstein, Dahai Li, Zhiyuan Liu, and	Narasimhan. 2024. τ -bench: A benchmark for tool-	772
716	Maosong Sun. 2024b. ToolLLM: Facilitating large	agent-user interaction in real-world domains. In <i>Ad-</i>	773
717	language models to master 16000+ real-world APIs.	<i>vances in Neural Information Processing Systems</i>	774
718	In <i>Proceedings of the 12th International Conference</i>	<i>(NeurIPS)</i> , volume 38.	775
719	<i>on Learning Representations (ICLR)</i> .	Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak	776
720	Aymeric Roucher, A Villanova del Moral, Thomas Wolf,	Shafraan, Karthik Narasimhan, and Yuan Cao. 2023.	777
721	Leandro von Werra, and Erik Kaunismäki. 2025.	ReAct: Synergizing reasoning and acting in language	778
722	smolagents: A smol library to build great agentic	models. In <i>Proceedings of the 11th International</i>	779
723	systems. <i>Hugging Face</i> .	<i>Conference on Learning Representations (ICLR)</i> .	780
724	Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta	Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga,	781
725	Raileanu, Maria Lomeli, Eric Hambro, Luke Zettle-	Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingn-	782
726	moyer, Nicola Cancedda, and Thomas Scialom. 2023.	ing Yao, Shanelle Roman, Zilin Zhang, and Dragomir	783
727	Toolformer: Language models can teach themselves	Radev. 2018. Spider: A large-scale human-labeled	784
728	to use tools. In <i>Advances in Neural Information Pro-</i>	dataset for complex and cross-domain semantic pars-	785
729	<i>cessing Systems (NeurIPS)</i> , volume 36.	ing and text-to-SQL task. In <i>Proceedings of the 2018</i>	786
730	Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté,	<i>Conference on Empirical Methods in Natural Lan-</i>	787
731	Yonatan Bisk, Adam Trischler, and Matthew	<i>guage Processing (EMNLP)</i> . Association for Compu-	788
732	Hausknecht. 2021. ALFWorld: Aligning text and	tational Linguistics.	789
733	embodied environments for interactive learning. In	Lingfeng Zeng, Fangqi Lou, Zixuan Wang, Jiajie Xu,	790
734	<i>Proceedings of the 9th International Conference on</i>	Jinyi Niu, Mengping Li, Yifan Dong, Qi Qi, Wei	791
735	<i>Learning Representations (ICLR)</i> .	Zhang, Ziwei Yang, Jun Han, Ruilun Feng, Ruiqi	792
736	Pat Verga, Sebastian Hofstatter, Sophia Althammer, Yix-	Hu, Lejie Zhang, Zhengbo Feng, Yicheng Ren, Xin	793
737	uan Su, Aleksandra Piktus, Arkady Arkhangorodsky,	Guo, Zhaowei Liu, Dongpo Cheng, and 2 others.	794
738	Minjie Xu, Naomi White, and Patrick Lewis. 2024.	2025. FinGAIA: A chinese benchmark for AI	795
739	Replacing judges with juries: Evaluating LLM gener-	agents in real-world financial domain. <i>arXiv preprint</i>	796
740	ations with a panel of diverse models. <i>arXiv preprint</i>	<i>arXiv:2507.17186</i> .	797
741	<i>arXiv:2404.18796</i> .	Yilun Zhao, Yunxiang Li, Chenying Li, and Rui Zhang.	798
742	Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa	2022. MultiHiertt: Numerical reasoning over multi	799
743	Liu, Noah A. Smith, Daniel Khashabi, and Hannaneh	hierarchical tabular and textual data. In <i>Proceedings</i>	800
744	Hajishirzi. 2023. Self-Instruct: Aligning language	<i>of the 60th Annual Meeting of the Association for</i>	801
745	models with self-generated instructions. In <i>Proceed-</i>	<i>ings of the 60th Annual Meeting of the Association for</i>	802
746	<i>ings of the 61st Annual Meeting of the Association</i>	<i>Computational Linguistics (ACL)</i> , pages 6588–6600.	803
747	<i>for Computational Linguistics (ACL)</i> , pages 13484–	Association for Computational Linguistics.	
748	13508. Association for Computational Linguistics.	Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan	804
749	Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten	Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin,	805
750	Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc V. Le,	Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang,	806
751	and Denny Zhou. 2022. Chain-of-thought prompt-	Joseph E. Gonzalez, and Ion Stoica. 2023. Judging	807
752	ing elicits reasoning in large language models. In	LLM-as-a-judge with MT-Bench and Chatbot Arena.	808
753	<i>Advances in Neural Information Processing Systems</i>	In <i>Advances in Neural Information Processing Sys-</i>	809
754	<i>(NeurIPS)</i> , volume 35.	<i>tems (NeurIPS), Datasets and Benchmarks Track</i> ,	810
755	xAI. 2025. Grok 3 beta: The age of reasoning agents.	volume 36.	811
756	Technical report, xAI.	Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou,	812
		Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue	813

Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Graham Neubig. 2024. WebArena: A realistic web environment for building autonomous agents. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*.

Fengbin Zhu, Wenqiang Lei, Youcheng Huang, Chao Wang, Shuo Zhang, Jiancheng Lv, Fuli Feng, and Tat-Seng Chua. 2021. TAT-QA: A question answering benchmark on a hybrid of tabular and textual content in finance. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (ACL-IJCNLP)*, pages 3277–3287. Association for Computational Linguistics.

A Cross-Benchmark Comparison of Size–Accuracy Trends

Beyond the within-benchmark sanity check shown in Figure 1, where accuracy on FinOpsBench grows log-linearly with the agent’s base-model size, we cross-check the size–accuracy relationship on two other recent financial-agent benchmarks. Parameter counts for closed-source models throughout this appendix come from the IKP calibration of Li (2026); we use the direct estimates from that paper’s main table where available and invert the published log-linear calibration on the extended-table IKP scores otherwise.

FinAgentBench (Choi et al., 2025). On the document- and chunk-ranking tasks of FinAgentBench (Figure 7), the trend is the opposite of the one we observe on FinOpsBench: across the three evaluated proprietary models (Claude-Sonnet-4 at ≈ 237 B, Claude-Opus-4 at ≈ 1.4 T, and GPT-o3 at ≈ 3 T), the *smallest* model wins or ties on every metric (nDCG@5, MAP@5, MRR@5) in both ranking tables. A monotonically negative slope on accuracy versus parameter count is a strong negative signal for benchmark validity. It suggests either that the task probes a narrow stylistic regime that the larger models are over-trained out of, or that smaller models in the same family have been exposed to closely related retrieval data during training. In short, performance on this benchmark appears to be governed more by model-family biases or potential data leakage than by raw model capability.

Finance Agent Benchmark (Bigeard et al., 2025). On the 22-model Finance Agent Benchmark leaderboard (Figure 8), accuracy does increase with parameter count, fitting a roughly +15.5 pp-per-decade log-linear trend. However, the residual variance around this trend is large enough that the fit is not, on its own, a reliable validity signal. First,




Figure 7: Size–accuracy on FinAgentBench (Choi et al., 2025), Tables 1 and 2. The smallest model wins or ties on every metric, an inverse trend that is inconsistent with what we observe on FinOpsBench.




Figure 8: Size–accuracy on the Finance Agent Benchmark (Bigeard et al., 2025), Table 2. The positive log-linear slope (≈ 15.5 pp/decade) is consistent with the FinOpsBench picture, but the wide intra-size scatter (20–40 pp) and the same-model reasoning-mode variation indicate that size is not the dominant driver.

at comparable sizes models differ by 20–40 pp; for example, Claude 3.7 Sonnet (≈ 676 B) reaches 44% while GPT-4o (≈ 720 B) reaches only 20%. Second, several large models sit far below the trend line: o1 at ≈ 3.5 T scores only 21.4%, and LLaMA 4 Maverick (Meta AI, 2025) at ≈ 400 B scores only 3.1%, which the source paper attributes to tool-interface misuse. Third, and most strikingly, the same model run in two different reasoning modes can differ by 13 pp at fixed parameter count (Grok 3 Mini (xAI, 2025) at 30.9% high-reasoning vs. 17.6% low-reasoning). The positive slope therefore conflates parametric capacity with reasoning paradigm and tool-use training; it does not isolate size as the dominant driver, in contrast to the cleaner relationship observed on FinOpsBench.

B Per-model accuracies across the two FinOpsBench versions

This appendix complements the side-by-side bar chart in Figure 4 with a per-model agreement view:




Figure 9: Per-model agreement between FinOpsBench-v1 and FinOpsBench-v2 accuracies. Models on the $y = x$ line score identically on both benchmarks; models below the line are weaker on FinOpsBench-v2.

Figure 9 plots the same data as a scatter against a $y = x$ reference line, separating models that score equally on the two versions from those that score lower on FinOpsBench-v2.

C Benchmark Statistics

This appendix collects descriptive statistics for both benchmark versions, organised by version below. Together the plots give a per-axis (text, schema, agent activity, tool surface) snapshot of how the two versions compare at the example level.

C.1 FinOpsBench-v1 Statistics

The figures below summarise FinOpsBench-v1 along three axes. Figures 10 and 11 cover the text side of the corpus (query length and the semantic spread of the queries in UMAP (McInnes et al., 2018) space). Figures 12 and 13 describe the size of the per-example backing stores in terms of table count and total row count. Figures 14 and 15 report the interaction depth of the reference agent traces (assistant turns and tool invocations per example), indicating how much agent activity each task induces.

The figures below summarise FinOpsBench-v2. Figures 16–18 cover the answer-length distribution, the system-prompt-length distribution, and the per-example tool count. Figures 19 and 20 project the system prompts and the tool-name vocabularies into UMAP space to visualise topical and lexical coverage.




Figure 10: FinOpsBench-v1: Query length distribution.




Figure 11: FinOpsBench-v1: UMAP of query embeddings.

D Example Queries from FinOpsBench-v1

The following queries are drawn at random from the filtered FinOpsBench-v1 test set. Each query is paired with the persona the LLM was instructed to assume when it generated the question; together they illustrate the breadth of user roles and the diversity of analysis categories covered by FinOpsBench-v1.

- *Internal Auditor*: Can you provide a list of transactions that were processed without obtaining necessary managerial approvals in the last quarter?
- *Tax Advisor*: Identify any vendor invoices that may impact VAT input tax claims for the current tax period.
- *Chief Financial Officer*: How do we ensure the reported accounts payable aging genuinely




Figure 12: FinOpsBench-v1: Number of data tables per example.




Figure 14: FinOpsBench-v1: Number of assistant turns per example.




Figure 13: FinOpsBench-v1: Number of total data rows per example.




Figure 15: FinOpsBench-v1: Number of tool calls per example.

reflects outstanding vendor obligations without distortion?

- *Financial Controller*: What is the total outstanding balance on all overdue invoices by supplier category?
- *Procurement Officer*: Which suppliers have outstanding invoices that have not been approved for payment yet?
- *Finance Director*: What percentage of supplier invoices in the last financial year were paid late, grouped by supplier?
- *Risk Manager*: List suppliers with overdue invoices over 90 days that have credit limit warnings.
- *Internal Controls Manager*: Detect any payments made with override flags but missing authorization memos.

Across the full 5,979-example set we observe more than 60 distinct user roles (e.g. CFO, internal/external auditor, controller, procurement officer, treasury manager, tax advisor, compliance officer, risk manager), and queries span the accounts-payable, reconciliation, variance-analysis, revenue-recognition, and financial-reporting categories described in Section 3.2.

E Example Questions from FinOpsBench-v2

The following questions are drawn at random from the filtered FinOpsBench-v2set. Each is a multi-hop numerical-reasoning question taken from a FinQA item; the agent must reach the answer by calling the example’s bespoke Python tools rather than by reading the original FinQA table directly.




Figure 16: FinOpsBench-v2: Answer length distribution.




Figure 18: FinOpsBench-v2: Number of tools per example.




Figure 17: FinOpsBench-v2: System prompt length distribution.




Figure 19: FinOpsBench-v2: UMAP of system prompt embeddings.

- What is the percent change in receivables from or (payables to) the Money Pool from 2001 to 2002?
- In 2018, what was the ratio of the qualified defined benefit pension plans' estimated payments for the period starting after 2024 compared with the amount for 2019?
- What percentage were North American consumer packaging net sales of consumer packaging sales in 2014?
- As of December 31, 2016, what was the percent of the total future minimum lease commitments for operating leases that was due in 2017?
- What was the change in percentage of the company's consolidated net sales from 2006 to 2008?
- For the year ended December 31, 2007, what was the ratio of the shares granted to the shares vested?
- If you held 1,000 shares on May 30, 2014, how much would you receive in dividends?
- What is the total, in millions, of current assets acquired?

Across the full 1,108-example set, every question is taken from FinQA and the surrounding tool

environment is generated automatically (see Section 3.3). The questions span ratio computations, year-over-year changes, percentage shares of totals, and aggregate sums, mirroring the financial-reasoning surface of the original FinQA dataset and reinstantiating it as an agentic task.

Table 5 shows a fully worked example: the question text, the bespoke Python tools the agent has access to, and the reference tool-call sequence that solves the task.

F Per-example file layout for FinOpsBench-v2

Each example in FinOpsBench-v2 is shipped as a small directory of source files, listed below.

`initial_solution.txt`

Correct answer to the question.

`synthetic_db_generator_augmented.py`

Python code that creates the SQLite database.




Figure 20: FinOpsBench-v2: UMAP of tool name embeddings.

Query (part of system prompt)
What is the percent change in receivables from or (payables to) the Money Pool from 2001 to 2002?
Available tools
<pre>list_available_years() get_quarterly_money_pool_amount(year, q) get_total_money_pool_amount_for_year(year) compute_percent_change(old, new) get_cash_flow_transactions(start, end, limit) sum_cash_flow_for_year(year)</pre>
Correct tool-call sequence
<pre>def annual(y): return sum(get_quarterly_money_pool_amount(y, q) for q in range(1, 5)) a01 = annual(2001) a02 = annual(2002) pct = compute_percent_change(a01, a02) answer = f"{round(pct)}%"</pre>

Table 5: A worked FinOpsBench-v2 example (excerpt): question, available tools, and reference tool-call sequence.

tools_augmented.py
Tools required for the multi-step processing, plus distractor tools.

tool_proxy.py
Same tools, adapted for use with smolagents.

correct_plan_augmented.py
Intended sequence of tool calls that solves the task.

agent_system_prompt.txt
LLM prompt with context, tool descriptions, and the question.

agent.py
Agent code that runs a specified LLM with the system prompt and tools.

In addition, each example directory carries auxiliary intermediate files kept from the build process

for reproducibility.

G Category distribution of FinOpsBench-v1

Table 6 reports the share of FinOpsBench-v1 examples assigned to each of the five categories from Section 3.2. Categorisation is done automatically by nearest-anchor TF-IDF similarity over a short description sentence per category, applied to the full FinOpsBench-v1 stream. The mix is dominated by accounts-payable queries (over half of the corpus), reflecting the seed-query composition; the remaining four categories are balanced between 7 and 15%.

Category	% of FinOpsBench-v1
Accounts Payable (AP) analysis	56.5
Variance analysis	14.8
Data integrity & reconciliation	13.5
Revenue recognition	8.2
Financial reporting	7.1

Table 6: Category share of the full FinOpsBench-v1 stream (percentages, computed by nearest-anchor TF-IDF assignment). Rounded to one decimal; totals to 100.1% due to rounding.

H Prompts and Examples

Figure 22 shows the prompt template used for the panel LLM judge in FinOpsBench-v1.

Figure 21 shows an annotated example from FinOpsBench-v1 illustrating the schema and data format.

I Use of AI Assistants

In the spirit of transparent reporting of AI-assisted authoring, we summarise how AI tools were used in the preparation of this paper. All core benchmark code, prompts, and configuration were written manually by the authors and reviewed end-to-end; the main body of the paper was likewise written manually. We used Claude (Anthropic, 2024) during the final preparation pass for two narrowly scoped tasks: (i) editorial revisions of phrasing and tightening of existing prose, and (ii) auxiliary code for collecting statistics about the benchmark and rendering some of the analysis figures. Every artifact produced by Claude was iter-

Query: Could duplicated invoice payments be causing discrepancies in our aged payables report?

Schema:

```
CREATE TABLE Suppliers (  
    supplier_id INTEGER PRIMARY KEY,  
    supplier_name TEXT NOT NULL  
);  
CREATE TABLE Invoices (  
    invoice_id INTEGER PRIMARY KEY,  
    supplier_id INTEGER NOT NULL,  
    invoice_date DATE NOT NULL,  
    due_date DATE NOT NULL,  
    invoice_amount NUMERIC NOT NULL,  
    FOREIGN KEY (supplier_id)  
    REFERENCES Suppliers(supplier_id)  
);  
CREATE TABLE Payments (  
    payment_id INTEGER PRIMARY KEY,  
    invoice_id INTEGER NOT NULL,  
    payment_date DATE NOT NULL,  
    payment_amount NUMERIC NOT NULL,  
    FOREIGN KEY (invoice_id)  
    REFERENCES Invoices(invoice_id)  
);
```

Data:

```
INSERT INTO Payments VALUES  
(5001, 1002, '2025-07-05', 500.00),  
(5002, 1002, '2025-07-10', 500.00),  
(5003, 1002, '2025-07-12', 500.00);
```

Figure 21: FinOpsBench-v1: Annotated example item showing query, schema, and data (excerpt).

Task

You are a senior QA engineer working on AI agents for finance. You have to judge the test case provided and check the following rules:

- The test case's data should be natural: it should be similar to what data could a real company have, and how it would be stored in a database.
- The agent's trace should be reasonable: the agent should look at the data needed to answer the query.
- The agent's trace should be sound: the reasoning steps should be logical, should follow from the previous steps and lead to the final answer.
- The agent's reasoning and final answer should be grounded: the only data the agent has access to is what was returned from its tool calls. The reasoning should not be based on any other data, including the data in the database that was not queried with the tools.
- The agent's final answer should answer the user's query.

In case one of these requirements are not met, you should provide a concise actionable feedback so the author can improve the test case.

Test case:

```
## Query  
{{ item.query }}  
  
## SQL schema  
{{ item.schema_sql }}  
  
## Data  
{{ item.data_sql }}  
  
## Agent trace  
{{ item.agent_dialog | tojson(2) }}  
  
# Output format  
Reply with JSON object like this:  
{  
    "data_is_natural": 1 | 0,  
    "trace_is_reasonable": 1 | 0,  
    "trace_is_sound": 1 | 0,  
    "reasoning_is_grounding": 1 | 0,  
    "answer_is_sound": 1 | 0,  
    "feedback": "..."  
}
```

Figure 22: FinOpsBench-v1: Prompt template for a panel member judge.

1067 actively reviewed and accepted by the authors; no
1068 Claude output was incorporated without manual
1069 inspection.